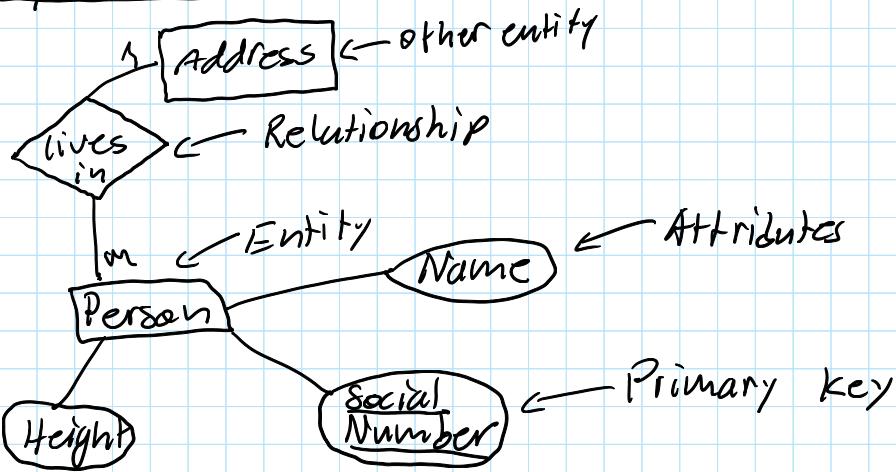


Entity Relationship Model



- Candidate keys: Set of attributes, that we can later choose one of to be the primary key

- Cardinality:

• 1:1

• 1:N

• N:1

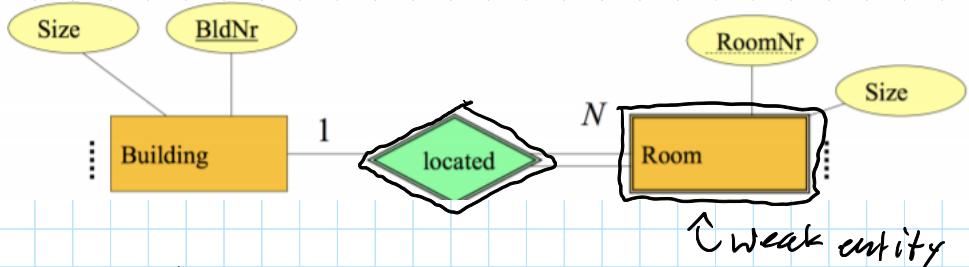
• N:M

} A relation, even of type 1:1, optionally connects entities!

- Weak entity:

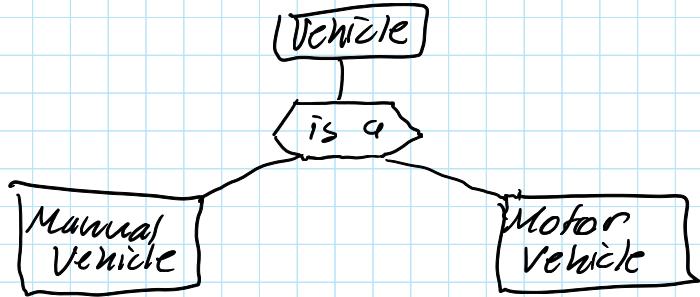
• Existence of room depends on the existence of the associated building

• Primary key of a room: BldNr + RoomNr



- Generalization:

• Like a parent



- A vehicle can be both manual-vehicle and motor-vehicle at the same time

- It does not require, that each entity of Vehicle is either a manual-vehicle or motor-vehicle

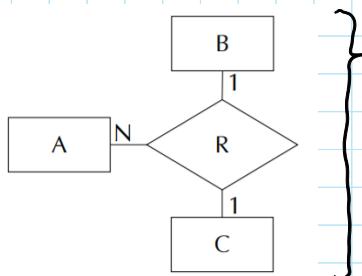
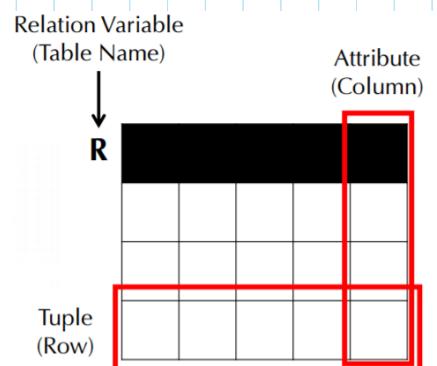
Keys

- Key: Must be unique
- Superkey: Any set of attributes, which identify an entry uniquely
- Candidate key: Minimal set of superkeys
- Primary key: Just one particular candidate key
- Foreigner key: Attribute in a table, which has a relation to another table, where that attribute is the primary key
 - For every foreign key one of the two conditions must hold:
 1. Value of foreign-key is NULL
 2. Referenced tuple must exist

Relation Model

- Student: $\{ \text{Legi: Integer}, \text{Name: String}, \text{Semester: Integer} \}$
- attends: $\{ \text{Legi: Integer}, \text{Nr: Integer} \}$

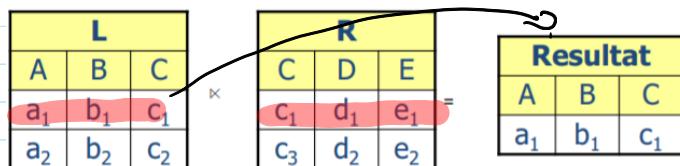
(Primary key müssen beide sein, da es sich um eine m:n-Relation handelt)



} Primary key: $R(\underline{A}, \underline{B}, \underline{C})$ or
 $R(\underline{A}, \underline{B}, \underline{C})$

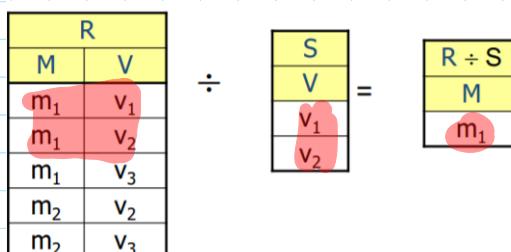
Relational Algebra

- Selection: $\sigma_{\text{semester} > 10}$ (student)
- Projection: Π_{level} (professor) : only outputs the different levels
- Cartesian product: Jede Zeile mit jeder anderen Zeile
- Natural join: $A \bowtie B$: Matches tuples (= inner join)
- Left outer join: $\bowtie L$: natural join + unmatched tuples from A
- Right outer join: $\bowtie R$ natural join + unmatched tuples from B
- Full outer join: \bowtie natural join + unmatched tuples from A/B
- Left semi join: \bowtie tuples from A matching tuples from B



- Right semi join: $\bowtie R$ tuples from B matching tuples from A
- Rename: Requirements \leftarrow prerequisite (requires)
- Intersection: $\cdot A \cap B$
 - Only works, if both tables have same attributes

- Relational division:



n-way join

- $A \bowtie B \bowtie C \Rightarrow 3\text{-way join}$
- Possible ways to order n-way join: $\frac{2^n}{n!}$

SQL (=structured query language)

Data types

- **character (n), char (n)** ← #chars
- **character varying (n), varchar (n)**
- **numeric (p,s), integer**
- **blob or raw** for large binaries
- **clob** for large string values (character large object)
- **date**

Create Tables

- **create table Professor**

```
(PersNr      integer not null,
 Name        varchar (30) not null
 Level       character (2) default „AP“);
```

Delete a Table

- **drop table Professor;**

Modify the structure of a Table

- **alter table Professor add column(age integer);**

Management of Indexes (Performance tuning)

- **create index myIndex on Professor(name, age);**
- **drop index myIndex;**

Delete tuples

```
delete Student
where Semester > 13;
```

Update tuples

```
update Student
set Semester = Semester + 1;
```

```
select PersNr, Name, Level
from Professor
order by Level desc, Name asc;
```

```
select distinct Level
```

```
from Professor
```

```
select Name
```

```
from Professor, Lecture
```

```
where PersNr = ProfNr and Title = 'Operating Systems';
```

$\prod \text{Name}(\sigma_{\text{PersNr} = \text{ProfNr} \wedge \text{Title} = 'Operating Systems'}(\text{Professor} \times \text{Lecture}))$

Select *
From A, B
WHERE A.id = B.id

≈

Select *
From A
Join B ON A.id = B.id

Select partid
FROM part

≈

Select p.partid
FROM part p

} Eliminate
duplicates

} From Rel. Algebra
to SQL

```

select s.Name, l.Title
from Student s, attends a, Lecture l
where s.Legi = a.Legi and
      a.Nr = l.Nr;
} shorter names

select Title, PersNr as ProfNr
from Lecture;
} Rename Attributes

(select Name
  from Assistant)
union
(select Name
  from Professor);
} union/intersect/minus

select avg (Semester)
  from Student;
} functions: avg, max, min, count, sum

select Name
  from Professor
where PersNr (not)in (select PersNr
  from Lecture);
} Subqueries
} where PersNr in (...) AND PersNr not in (...)
```

Null-Values in SQL

- $0 + \text{null} \Rightarrow \text{null}$
- $0 * \text{null} \Rightarrow \text{null}$
- $\text{null} = \text{null} \Rightarrow \text{unknown}$
- $\text{null} < 13 \Rightarrow \text{unknown}$

Cases

```

select Legi, (
  case when Grade >= 5.5 then 'sehr gut'
    when Grade >= 5.0 then 'gut'
    when Grade >= 4.5 then 'befriedigend'
    when Grade >= 4.0 then 'ausreichend'
    else 'nicht bestanden' end)
from tests;
```

Group by / having
group by origin
having count (order id) < 10

Comparisons

- "%" represents any sequence of characters (0 to n)
- "_" represents exactly one character
- N.B.: For comparisons with = , % and _ are normal chars.

```
select *  
from Student  
where Name like 'Kossman%';
```

Views

```
CREATE VIEW FlightFromPEK AS  
SELECT * FROM Flight WHERE orig=PEK;
```

```
SELECT *  
FROM FlightFromPEK  
WHERE dest = ZRH;
```

} inline
With FlightFromPEK AS (-)

- Not updatable :
 - MAX/COUNT/MIN/AVG
 - DISTINCT - keyword
 - JOIN
 - GROUP BY or HAVING
 - More than one table
 - Subqueries
 - Not include the primary key in select
 - Arithmetic expressions

Recursion

With

AncestorOfD as (

```
Select parent  
from parentof  
where D = child;
```

UNION

```
Select p2.parent  
FROM AncestorOfD,  
      ParentOf  
WHERE AncestorOfD.ancestor = parentof.child;
```

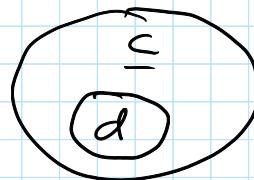
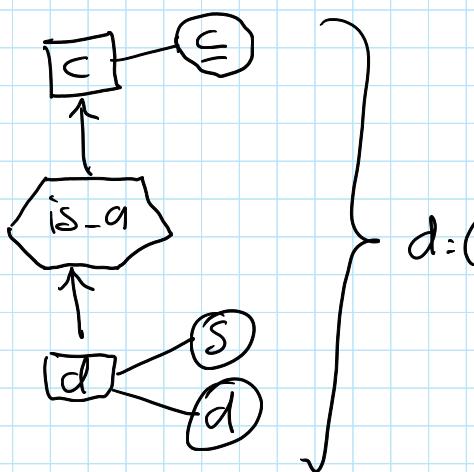
Recursive call of

Converting Relationships



- $M:N \Rightarrow A:(\underline{a}) / B:(\underline{b}) / \text{rel}:(\underline{a}, \underline{b})$
- $1:N \Rightarrow A:(\underline{a}) / B:(\underline{b}, a)$
- $1:1 \Rightarrow \text{rel}:(\underline{a}, \underline{b}) \text{ or } \text{rel}:(a, \underline{b})$

Specification



$\hat{c}d$ is a subset, if we know c , we also know d

Constraints

- How database makes sure changes are consistent and do not cause trouble
- Problems:
 - Inserting tuples without a key
 - Adding references to tuples that do not exist
 - Negative age
- CREATE TABLE Student


```

      (sID int primary key,
      sName text, constraint
      GPA real not null);
      
```

constraint

- Only one primary-key, but if some other column is as well unique, use keyword "unique"

- Primary key over two columns:

```
CREATE TABLE Student
```

```
(sID int,  
sName text,  
GPA real,  
primary key (sID, sName));
```

- Primary key can't be NULL

- Unique attribute can be NULL

- Checking constraints

- GPA real check [$(GPA \leq 4.0 \text{ and } GPA > 0.0)$,
or (...)]

- foreign key : FOREIGN KEY (PersonID) REFERENCES Persons(PersonID)

- on delete cascade

↳ If record is deleted in Persons,
it is also deleted in the
new table

- on delete set null

↳ If record is deleted in Persons,
its foreign-key entry in the new
table is set to null

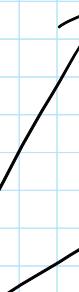
- on delete set default

- on delete no action
- on delete restrict

↳ Same for on modify

constraint key
primary key (...)

constraint test
check (...)



Redundancy

- Update anomaly: We need to update the general information in every row
- Insertion anomaly: Add a new general element row
- Deletion anomaly: If we remove all elements, the general element is also deleted

Functional dependencies

Professor: {[PersNr, Name, Level, Room, City, Address, Zip, AreaCode, Canton, Population, Direktion]}

- > $\{ \text{PersNr} \} \rightarrow \{ \text{PersNr, Name, Level, Room, City, Address, Zip, AreaCode, Canton, Population, Direktion} \}$
- > $\{ \text{City, Canton} \} \rightarrow \{ \text{Population, AreaCode} \}$
- > $\{ \text{Zip} \} \rightarrow \{ \text{Canton, City, Population} \}$
- > $\{ \text{Canton, City, Address} \} \rightarrow \{ \text{Zip} \}$
- > $\{ \text{Canton} \} \rightarrow \{ \text{Direktion} \}$
- > $\{ \text{Room} \} \rightarrow \{ \text{PersNr} \}$

- $\alpha \subseteq R$ is a superkey if $\alpha \rightarrow R$

Minimal Basis

- Minimal set of functional dependencies

1. Left side reduction

$$\cdot \overbrace{xy}^{\subseteq} \rightarrow A$$

$$\cdot y \rightarrow x$$

$$\hookrightarrow y \rightarrow A$$

2. Right side reduction

$$\cdot CF \rightarrow BD$$

$$\cdot BC \rightarrow D$$

$$\hookrightarrow CF \rightarrow B$$

3. Eliminate FDs: $x \rightarrow \emptyset$

4. Apply union rule to FDs with same left side

Decomposition of relations

- $R = R_1 \cup R_2$

↳ Only proves, that relation is lossless and not, that it is lossy!

- Decomposition is lossless, if

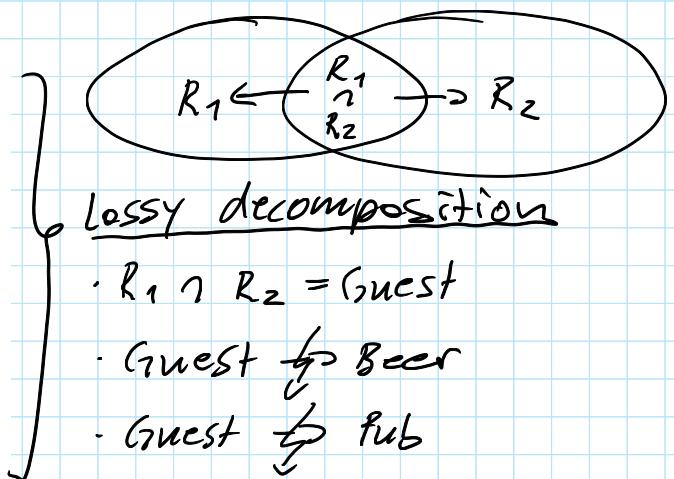
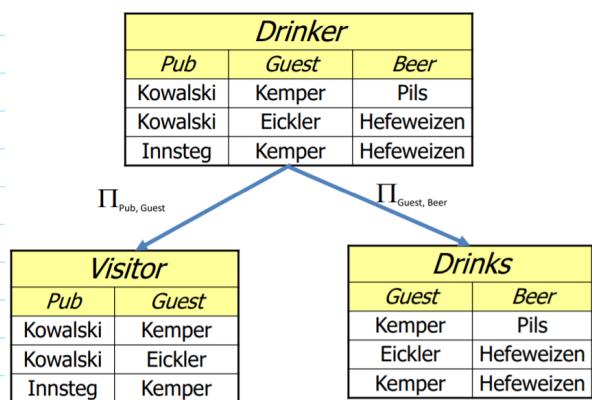
• $(R_1 \cap R_2) \rightarrow R_1$ or

↳ $S = S_1 \cup S_2$

• $(R_1 \cap R_2) \rightarrow R_2$

$S = S_1 \bowtie S_2$

Intersection



Normal Forms / Normalization

- Technique of organizing the data into multiple related tables to **minimize** data-redundancy

1st normal form

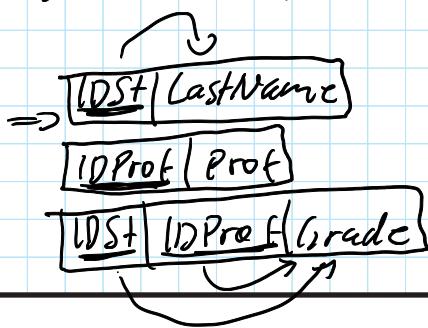
- Each attribute should only contain one element
- A column should contain values of same type
- Each column/attribute should have a unique name

2nd normal form (= Dependency on primary key)

- All other columns must fully depend on all primary-keys (not just one primary key)

A diagram showing a 3x5 table being decomposed into three separate relations. The table has columns labeled IDSt, LastName, IDProf, Prof, and Grade.

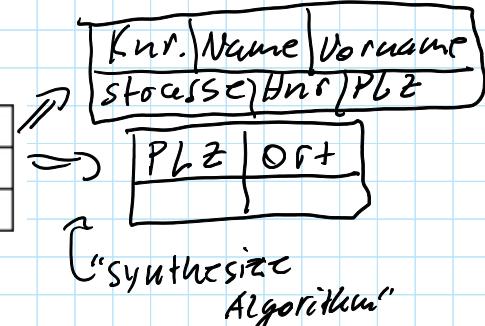
IDSt	LastName	IDProf	Prof	Grade
1	Mueller	3	Schmid	5
2	Meier	2	Borner	4
3	Tobler	1	Bernasconi	6



3rd Normal Form (=No transitive Dependencies)

- If a attribute is dependant on a non-primary-key attribute

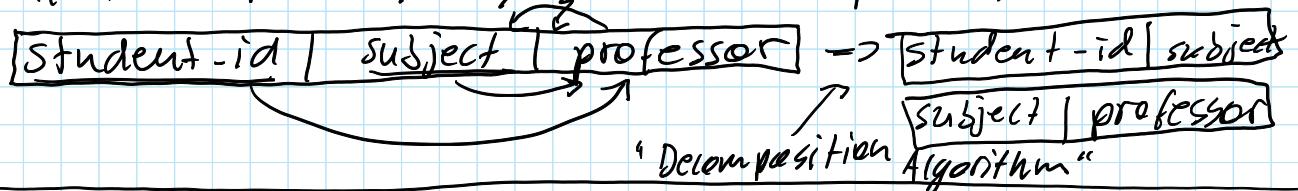
Kunde						
Knr.	Name	Vorname	Straße	Hnr.	PLZ	Ort
007	Mustermann	Max	Musterstr.	1	12345	Musterort



- Algorithm:
 1. compute minimal Basis
 2. For each func. dependency create table
 3. If needed, add table only containing old primary keys

BCNF / 3.5 Normal form

- If a non primary-key derives a primary-key

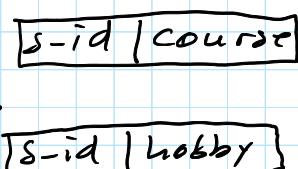


4th normal form

- No multi value dependencies (no 1:n or m:n relations)
- Multi-value-dependency

- If for a value single value of A, more than one value of B exists
- Have at least 3 columns

s_id	course	hobby
1	Science	Cricket
1	Maths	Hockey



- s_id → course
- s_id → hobby } student can have multiple courses/hobbies

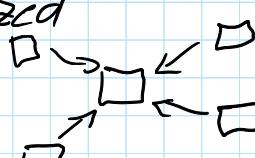
OLTP

- Online transaction processing
- Workload with lots of updates
- Run over 3NF and many tables
- E.g. E-commerce

OLAP

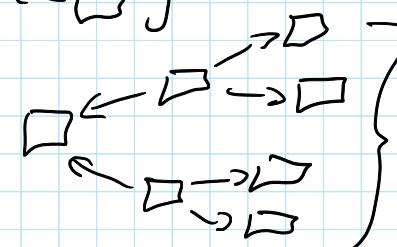
- Online analytical processing
- Workload with heavy and complex queries
- Take hours to process
- Run denormalized

- Star schema:



The diagram shows a central rectangular box representing the fact table, with four arrows pointing from it to four smaller rectangular boxes arranged around it, representing dimension tables. A curly brace to the right of the dimension tables is labeled "Denormalized".

- Snowflake schema:

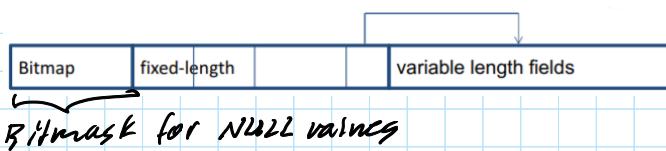


The diagram shows a central rectangular box representing the fact table, with four arrows pointing from it to four smaller rectangular boxes arranged around it. These smaller boxes then have their own internal connections. A curly brace to the right of the outer dimension tables is labeled "Normalized form of star schema".

- E.g. Weather forecast

How data is stored

- Records: Tuples
- Pages: collections of records from same table
- Blocks: Parts / collections of pages
- Tablespace: space for a database on a disk
- Structure of a record

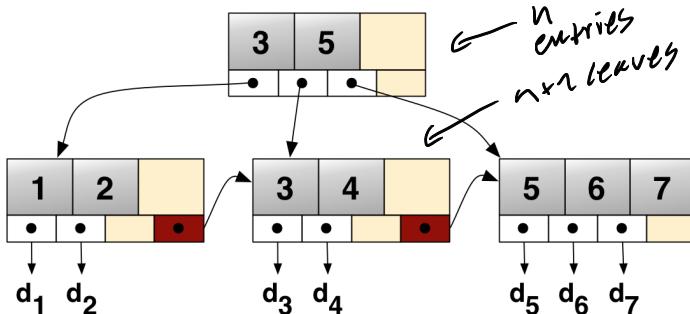


Structure of a page

- Header at the end of the page
 - Number of records in the current table
 - Slot directory: 0: slot free
1: slot used
- Unpacked Order: Free/used slots not ordered
- Packed order: Free/used slots ordered

Query Algorithms

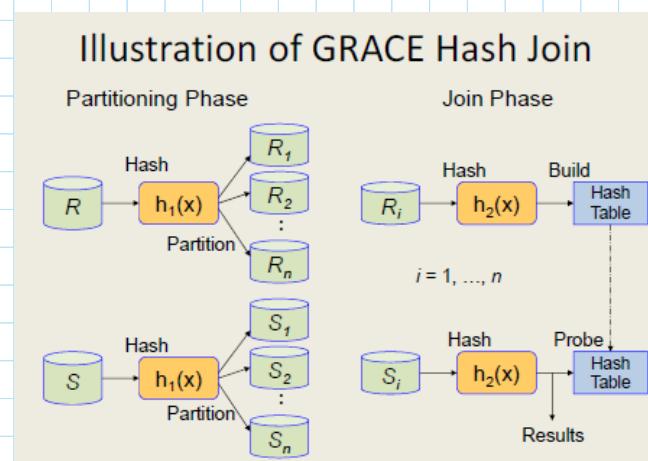
- Query selectivity: #Tuples in result vs #Tuples in DB
- Attribute cardinality: How many distinct values for attribute
- Scan
 - Iterate over every tuple in the DB
- Indexes
 - Avoid a full table scan \Rightarrow Build an index over some attributes
 - B+ tree:
 - Self balancing tree
 - Each node can have more than n children
 - Search, Insertion, Deletion = $O(\log(n))$



- Two ways for storing data:
 - pointer to tuple
 - tuple inside tree
- Hash-index
 - Every item is read/written once
- External Sort (= One pass sort)
 - Data does not fit in main-memory
 - N: # pages
 - M: # pages that can be stored in the buffer
- 1. Phase: Load tuples in buffer, sort, write back,
(create runs) do this for all tuples
- 2. Phase: Load from each run the x-th tuple
(merge runs) and do sorting again

- Joins

- Nested loop join: $O(A \cdot B)$
- Sort-merge join: $O(A \cdot \log(A) + B \cdot \log(B))$
- Canonical hash join:
 1. Hash table A
 2. Match table B with same hash function $\cdot O(A + B)$
- Grace Hash join:
 - If we have too many entries for one hash function



$\cdot O(A + B)$

Query Execution

1. Query is translated into relational-algebra-tree
 2. Tree is converted into different plans
 3. Cost of each plan is approximated
 4. Cheapest plan is chosen
- Make different trees

$$\begin{aligned} \cdot \sigma_{\theta_1 \wedge \theta_2} &= \sigma_{\theta_1}(\sigma_{\theta_2}(E)) \\ \cdot \sigma_{\theta_1}(\sigma_{\theta_2}(E)) &= \sigma_{\theta_2}(\sigma_{\theta_1}(E)) \\ \cdot \pi_{t_1}(\pi_{t_2}(\pi_{t_3}(E))) &= \pi_{t_1}(E) \\ \cdot E_1 \bowtie E_2 &= E_2 \bowtie E_1 \\ \cdot (E_1 \bowtie E_2) \bowtie E_3 &= E_1 \bowtie (E_2 \bowtie E_3) \\ \cdot \sigma_{\theta}(E_1 \bowtie E_2) &= [\sigma_{\theta}(E_1)] \bowtie E_2 \\ \cdot E_1 \cap E_2 &= E_2 \cap E_1 \\ \cdot E_1 \cup E_2 &= E_2 \cup E_1 \\ \cdot \sigma_{\theta}(E_1 - E_2) &= \sigma_{\theta}(E_1) - E_2 \end{aligned}$$

- The earlier we process selections, less tuples we need to manipulate higher up in the tree

Transaction

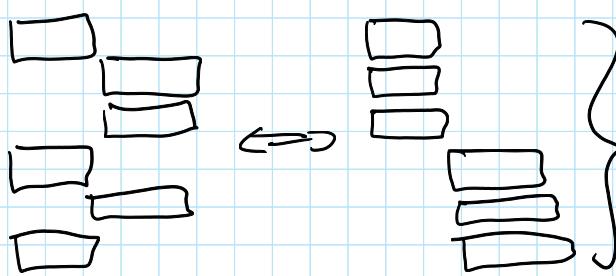
- Series of commands, extracted from a query
- Defines a transition between two different states
- Ends in commit (=finished) or abort (=rollback)
- ACID principle

· Atomicity: a transaction is executed in its entirety or not at all

- Consistency:  consistent state → transaction → consistent state
- Isolation: a transaction executes as if it were alone in the system ↳ 2-phase locking
- Durability: committed changes of a transaction are never lost and can be recovered ↳ Recovery

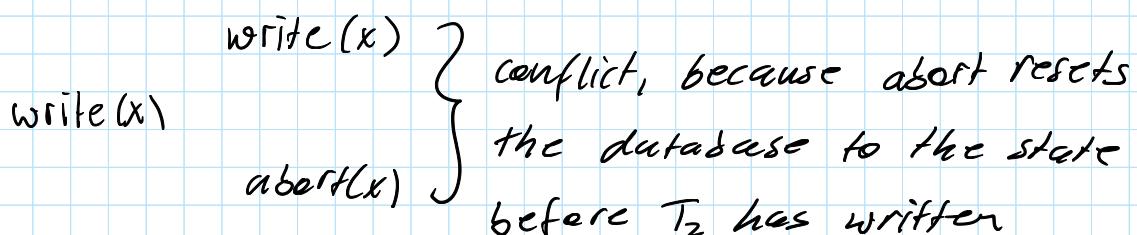
- Concurrency: conflict serializable

- We can translate two transactions into a serial schedule with a sequence of non-conflicting swaps
non-conflicting swaps
Reordering allowed

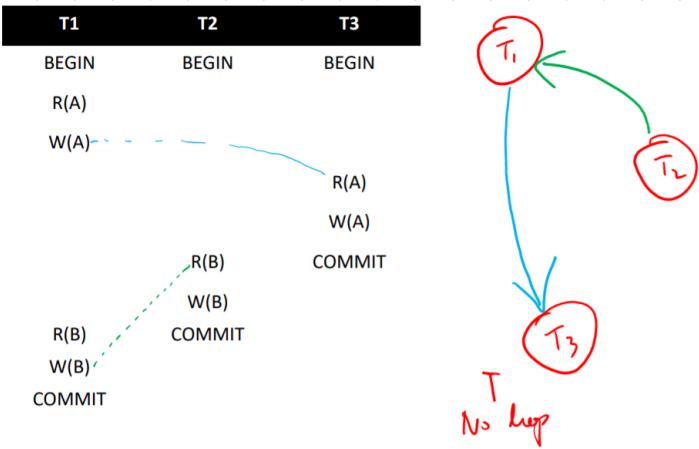


Wenn wir es so transformieren können mit non-conflict swaps, ist Schedule conflict-serializable

- conflict with aborts



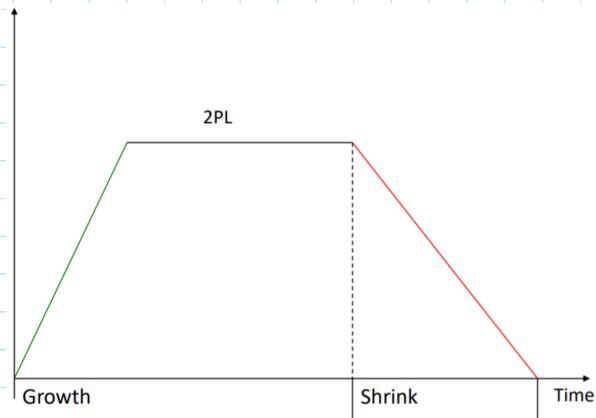
- Conflict serializable, iff dependency-graph is acyclic



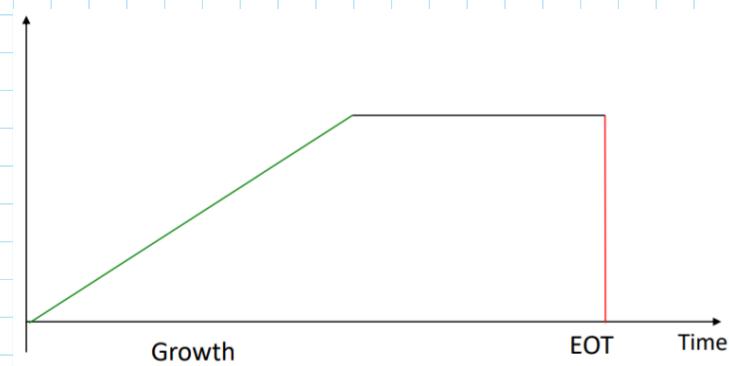
- Pessimistic synchronization

- Share lock: Only allows read operations
- Exclusive lock: Need for write operations

- Z-phase locking



- strict z-phase locking



- Optimistic synchronisation

- Snapshot isolation

Recovery

1. Application, System

- R1 Recovery: Undo a single transaction

2. System crash: lose main memory, keep disk

- R2 Recovery: Redo committed Transactions

- R3 Recovery: Undo active Transactions

3. System crash: loss of disks

- R4 recovery: Read backup of DB

- Recoverable:

Recoverable [edit]

Transactions commit only after all transactions whose changes they read, commit.

$$F = \begin{bmatrix} T1 & T2 \\ R(A) & \\ W(A) & \\ & R(A) \\ & W(A) \\ Com. & \\ & Com. \end{bmatrix} F2 = \begin{bmatrix} T1 & T2 \\ R(A) & \\ W(A) & \\ & R(A) \\ & W(A) \\ Abort & \\ & Abort \end{bmatrix}$$

If T_2 reads from T_1 , T_1 must commit first

These schedules are recoverable. F is recoverable because T1 commits before T2, that makes the value read by T2 correct. Then T2 can commit itself. In F2, if T1 aborted, T2 has to abort because the value of A it read is incorrect. In both cases, the database is left in a consistent state.

Avoids cascading aborts / rollbacks (ACA) [edit]

Also named cascadeless. Avoids that a single transaction abort leads to a series of transaction rollbacks. A strategy to prevent cascading aborts is to disallow a transaction from reading uncommitted changes from another transaction in the same schedule.

The following examples are the same as the ones in the discussion on recoverable:

$$F = \begin{bmatrix} T1 & T2 \\ R(A) & \\ W(A) & \\ & R(A) \\ & W(A) \\ Com. & \\ & Com. \end{bmatrix} F2 = \begin{bmatrix} T1 & T2 \\ R(A) & \\ W(A) & \\ & R(A) \\ & W(A) \\ Abort & \\ & Abort \end{bmatrix}$$

In this example, although F2 is recoverable, it does not avoid cascading aborts. It can be seen that if T1 aborts, T2 will have to be aborted too in order to maintain the correctness of the schedule as T2 has already read the uncommitted value written by T1.

The following is a recoverable schedule which avoids cascading abort. Note, however, that the update of A by T1 is always lost (since T1 is aborted).

$$F3 = \begin{bmatrix} T1 & T2 \\ R(A) & R(A) \\ W(A) & \\ & W(A) \\ Abort & \\ & Commit \end{bmatrix}$$

Note that this Schedule would not be serializable if T1 would be committed. Cascading aborts avoidance is sufficient but not necessary for a schedule to be recoverable.

- strict history: A schedule in which a transaction can neither **read or overwrite** an item x until the last transaction that wrote x has committed

On a cascadeless schedule a transaction T_2 cannot read a value a if a transaction T_1 wrote a before that and didn't commit. On a strict schedule T_2 also wouldn't be able to write a after T_1 wrote it (even if it read a before T_1 wrote it).

If you read carefully, the definition of strict says "not read **or overwritten**". That's the difference.

- Undo-logging: (= Rückgängig machen)

- Some transactions are in the middle, how to undo them?

· Start T
old value

· Update $\langle T, X, v \rangle$

· Commit T / Abort T

· If commit is in the log, it does not mean it is on the disk

· If T modifies database element X , write $\langle T, X, v \rangle$ to disk, before change of X is written to disk
old value

· If a transaction commits, commit must be written to disk only after all other changes are written to disk

- Redo logging (= Wiederholen)

· Update $\langle T, X, v \rangle$
new value

- Undo/Redo Logging

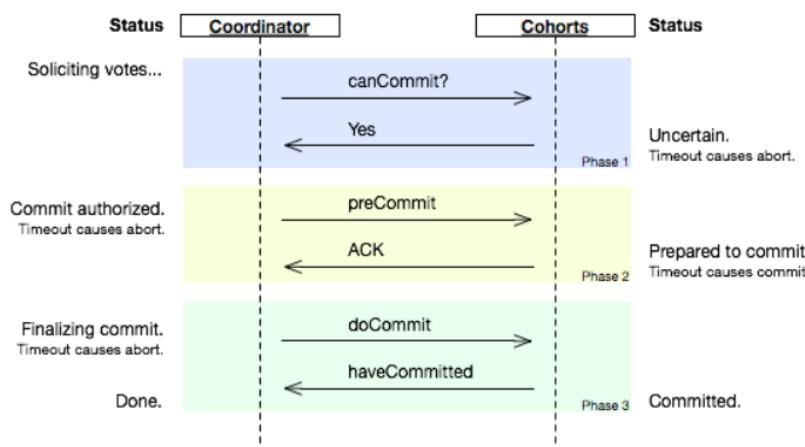
· Update $\langle T, X, v, w \rangle$ (= Store both, old and new value)

· Flush log, whenever Commit appear in log

Consensus (= Verteiltes Commit Protokoll für DB-Cluster)

- Problem of reaching an agreement among all working processes on the value of a variable
- 2-phase commit (in most cases the agent who will alter data)
 1. Coordinator sends Vote-Request to all participants
 2. All participants sends back Vote
 3. Coordinator collects all votes
 - 3.1. All Yes: Commit and send commit to all others
 - 3.2. Some No: Abort and send Abort to all which voted yes
- If time-out for waiting for response \rightarrow Abort
- If time-out for waiting for Vote-Req \rightarrow Abort local
- If time-out waiting for a decision \rightarrow uncertainty period
- Problem that leads to inconsistency:
 1. Coordinator sends Commit to first Agent
 2. Agent commits
 3. Coordinator and first Agent fails

3-phase commit



} introduced by
3-phase commit

Data Replication

- When are updates propagated?
 - Synchronous (No inconsistencies / worse response time)
 - Asynchronous (Good response time / Data inconsistencies)
- Where the update can be made?
 - Primary copy \Rightarrow master (Load at primary-copy is large)
 - Update everywhere \Rightarrow group
 - \hookrightarrow (copies need to be synchronized)

Key-Value-Store

- Map with keys and associating values
- Clients can put and request values per key