

# Systems Programming and Computer Architecture

C

- Static: Value bleibt unverändert zwischen Calls

- char: 1 byte

- short: 2 bytes

- int: 4 bytes

- long: 8 bytes

 Integer

- float: 4 bytes

- double: 8 bytes

 Floating point numbers

- Boolean: · 0: false

- !0: true (anything other than 0)

- Casting: the bits don't change, only the interpretation

- Arrays: C-compiler does not check array bounds

- Strings: · C has no real string-type

- Array of chars, terminated with null "10"

- char str[6] = { "h", "e", "l", "l", "o", "10" };

- char str[6] = "hello";

- Bit-level-operator: · And & / · OR | / · XOR: ^ / · NOT ~

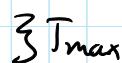
- Logic-operator: · And && / · OR || / · Not!  sign bit

- Arithmetic shift: If it's a signed integer, MSB is replicated back

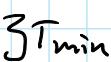
- Logical shift: If it's an unsigned integer, 0 are filled

## Numbers

- Two's complement: · Most significant bit = sign bit: 1=negativ / 0=positive

- größte Zahl = 011...1  Tmax

- zero = 00...0

- kleinste Zahl = 100...0  Tmin

- Reverse: 1. invert all bits

- 2. Add 1 to the right (=lsb)

- Unsigned numbers:
  - größte Zahl: 11...1 } 1111
  - kleinste Zahl: 00...0 } 0000
- Mixing signed/unsigned:
  - Signed  $\rightarrow$  Unsigned
  - Unsigned  $\rightarrow$  Unsigned

$\hookrightarrow -1 > 0 \text{ unsigned} \Rightarrow \text{true}$  ( $-1_{\text{signed}} \rightarrow 100\dots0 \rightarrow \text{Big unsigned}$ )
- Unsigned divide:  $x / 2^k = x \gg k$
- Signed divide:
  - $x / 2^k = x \gg k$  for  $x > 0$
  - $x / 2^k = x + (2^k - 1) \gg k$  for  $x < 0$  } Arithmetic shift
- Overflow:
 

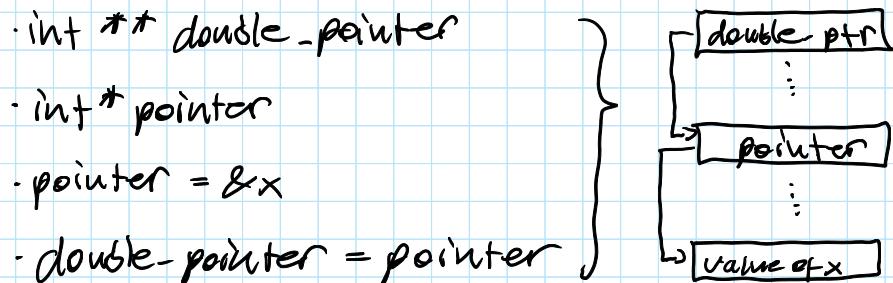
$\hookrightarrow$  We forgot the carry bit and treat the remaining bits as a normal number  
 $\hookrightarrow$  Causes over/underflow

## Pointers

- Pointers correspond to memory-addresses
- Functions have pointers as well
- $\&x$ : address    $\&x$ : value at address
- $\text{type} * \text{name};$  (points somewhere in the memory space)
 

$\hookrightarrow$  (illegal) to dereference it, because it points to NULL (0x00..0)
- $\text{type} * \text{name}$  // declares a pointer
- $\text{int test} = 5$
- $\text{name} = \&\text{test}$
- $*\text{test} = 10$

- Array: Arrays are treated as a pointer to the first element ( $A[i] \equiv *(A+i)$ )
- Size: `sizeof(type)` or `sizeof(value)`
- Double pointer:
  - Pointer to a pointer
  - `int ** double_pointer`
  - `int * pointer`
  - `pointer = &x`
  - `double_pointer = pointer`
- int-pointer zeigt auf word line (Byte-wise access)
- char-pointer zeigt auf erstes Bit (Bit-wise access)



### Passing values

- Passing by reference: Normally only objects
- Passing by value: primitives like int
- In C, we can use pointers to get a pseudo pass by ref.
  - ↳ Callee still receives a copy, but now it's a pointer
- A pointer can point to NULL, a reference not

### Memory

- Global variables:
  - allocated when program is run
  - deallocated when program exits
- Local variables:
  - allocated when method is called
  - deallocated when method exits
- C has no garbage collector: explicitly allocate/deallocate memory
- `malloc`:
  - `void* malloc(size_t)`
  - Returns a pointer to first byte of memory
  - Returns NULL, if memory cannot be allocated

- **calloc**: `void* calloc(size num, size sz)`
    - Allocates a block of memory with size  $num \times sz$
    - Memory is filled with 0
  - **free**: `void free(void*)`
    - Releases memory at the pointer
    - Must point to the first byte of allocated memory
  - **realloc**: `void* realloc(void* ptr, size)`
    - Changes the size of the pointer
    - Copies the data to a new location
- 

## Heap / Stack

- **Stack**: variables without malloc allocation
  - **Heap**: allocated by malloc
  - **Memory leak**: When code doesn't deallocate memory
- 

## Struct

- Like a class in Java, but without methods
  - `typedef struct quadra{`

```

    int length;
    int height;
}
```

  - `quad quadra1`  $\Leftarrow$  struct quadra quadra1
  - Use `"."` to refer to fields in a struct
  - Use `"->"` to refer to fields through a pointer to a struct  $\Rightarrow \cong (*p).x$
- 

## Generic pointer

- `void *generic_ptr`
- Can be converted to any type of pointer without casting
- `int *test = generic_ptr`

## Union

- Like a struct, but holds only one field at a time
- Union u $\in$   
int val;  
float fval;  
3

## The C-Preprocessor (text substitution tool)

- Include source code:
  - `#include <system-header.h>`
  - `#include "own_header.h"`
- Header-file:
  - file.h<sup>↑</sup> Everything with # is for the preprocessor
  - contains only declarations

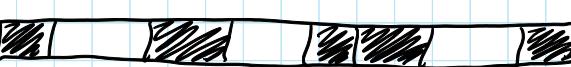
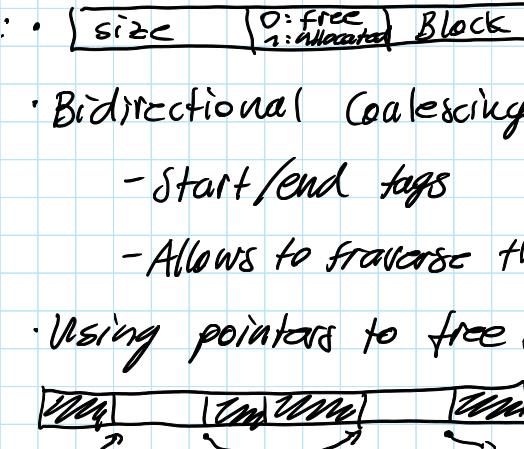
## Jumps

- `setjmp(env)`
  - saves the current state in env
  - Returns 0
- `longjmp(env, val)`
  - causes another return to the point saved by env
  - This return returns val
- (nut für Coroutines (siehe Grafik) oder Error-Handling)

```
void routineA()
{
    int r;
    printf("(A1)\n");
    r = setjmp(bufferA);
    if (r == 0) routineB();
    printf("(A2) r=%d\n", r);
    r = setjmp(bufferA);
    if (r == 0) longjmp(bufferB, 10001);
    printf("(A3) r=%d\n", r);
    r = setjmp(bufferA);
    if (r == 0) longjmp(bufferB, 20002);
    printf("(A4) r=%d\n", r);
}
```

```
void routineB()
{
    int r;
    printf("(B1)\n");
    r = setjmp(bufferB);
    if (r == 0) longjmp(bufferA, 10001);
    printf("(B2) r=%d\n", r);
    r = setjmp(bufferB);
    if (r == 0) longjmp(bufferA, 10002);
    printf("(B3) r=%d\n", r);
    r = setjmp(bufferB);
    if (r == 0) longjmp(bufferA, 10003);
}
```

## Memory Allocation

- Used by malloc/calloc/realloc / free
- Current heap size: wird während Programmablaufes  
immer größer
- Peak memory utilization:
  - Am besten 100% (keine Verschwendungen)
  - Wie viel Prozent des heaps ist  
$$\frac{\text{Maximum summe allocated Memory}}{\text{Heap size}}$$
- Fragmentation:
  - Internal: 
  - External: 
- Implicit list:
  - 
  - Bidirectional Coalescing of free blocks
    - Start/end tags
    - Allows to traverse the list backwards
- Explicit list:
  - Using pointers to free blocks
- Segregated free list:
  - Different free lists for different block sizes
  - For larger sizes: one list for each power of two

## Garbage Collector

- Common in functional languages, e.g. Java
- We can identify garbage, if there are no pointers to them
- Mark and sweep garbage-collector with malloc/free
  - 1. Use extra mark-bit in the head of each block
  - 2. Start at root\* and set mark bit on each reachable block
  - 3. Scan all blocks, and free them, if they are not marked

\*e.g. global variables

## Assembly

- Complex instruction set >< Reduced instruction set
- Transfer data between memory and registers
- Operand types:
  - Immediate: integer
  - Register:  $\%eax$
  - Memory: memory address given by register  
 $\hookrightarrow (\%eax)$
- Memory addressing:  $D(R_b, R_i, s) \hookrightarrow Mem[Reg[R_b] + s * Reg[R_i] + i]$   
Constant Base Index Scale Register Register
- Condition codes
  - Single bit registers: carry flag / zero flag / sign flag / overflow flag
  - Implicitly set by arithmetic-operations
  - Explicitly set by:
    - `cmplx Src2, Src1` (Arithmetic)
    - `testx Src2, Src1` (Bitwise)
- Conditional branch

1. Setup
  2. Test branch
  3. Body 1 if branch is taken
  4. Body 2 else branch is taken
  5. Finish
- 

## Linker Symbols

- Global symbol:
  - Defined in module, that can be referenced by other
- External symbol:
  - Global symbol that are referenced by module m, but defined in module n
- Local symbol:
  - Symbol that is initialized and only used in the same module (static int temp = 5)
  - but temp is not a symbol

- Strong symbol:
    - procedures and initialized values
  - Weak symbol:
    - Uninitialized global
    - compiler doesn't know, if it's an external variable, or an uninitialized one
  - Only one strong symbol
  - One strong symbol + several weak symbols } With the same name
  - Several weak symbols, pick random one
  - Use prefix extern, if you "import" the variable
- 

### Libraries

- Static libraries: 
  - Shared/dynamic libraries:
    - object files, that are loaded into application at run-/compile-time
- 

### Executable and linkable format

- ELF
- Anything that ends in .o is an ELF-file
- .text section: Code
- .data section: initialized global variables
- .bss:
  - Bullshit section
  - Uninitialized global variables

## Floating Point

- $\underbrace{(-1)^S}_{\text{Sign}} \cdot \underbrace{M}_{\text{Fraction}} \cdot \underbrace{2^E}_{\text{Exp}} \stackrel{?}{=} [\text{sign} | \text{frac}]$
- Single precision: 32 Bit: float
- Double precision: 64 Bits: double
- Casting: floating point  $\rightarrow$  int: rounding to zero: 4,56  $\rightarrow$  4
- Normalized values (good for big values)
  - Condition:  $\text{exp} \neq 0\dots00$  or  $\text{exp} \neq 1\dots11$
  - Exp: unsigned (coded as biased)
  - Bias:  $2^{e-1}-1$ , where  $e$  is the number of exponent bits
  - $E = \text{Exp} - \text{Bias}$
  - Fraction  $M = 1. xxxx\dots x_2$ 
    - Minimal:  $M = 000\dots0 \stackrel{?}{=} M = 1.0$
    - Maximal:  $M = 111\dots1 \stackrel{?}{=} M = 2.0$
- Denormalized values (good for small values near zero)
  - Condition:  $\text{exp} = 00\dots0$
  - Exponent value:  $E = -\text{Bias} + 1$
  - $\text{exp} = 00\dots0, \text{frac} = 000\dots0$ : represents value 0
  - $\text{exp} = 00\dots0, \text{frac} \neq 000\dots0$ : Number very close to 0.0
  - Fraction  $M = 0.xxxx\dots x_2$
- Special values
  - Condition:  $\text{exp} = 11\dots1$
  - $\text{frac} = 000\dots0$ :  $\pm\infty$
  - $\text{frac} \neq 000\dots0$ : not a number (=NaN)
- Rounding: Nearest even: if it's in the middle, round to even
  - $1,50 \rightarrow 2$
  - $2,50 \rightarrow 2$

	0 0000 000	-6	0	
	0 0000 001	-6	$1/8 * 1/64 = 1/512$	closest to zero
Denormalized numbers	0 0000 010	-6	$2/8 * 1/64 = 2/512$	
	...			
	0 0000 110	-6	$6/8 * 1/64 = 6/512$	
	0 0000 111	-6	$7/8 * 1/64 = 7/512$	largest denorm
	0 0001 000	-6	$8/8 * 1/64 = 8/512$	smallest norm
	0 0001 001	-6	$9/8 * 1/64 = 9/512$	
	...			
	0 0110 110	-1	$14/8 * 1/2 = 14/16$	
	0 0110 111	-1	$15/8 * 1/2 = 15/16$	closest to 1 below
Normalized numbers	0 0111 000	0	$8/8 * 1 = 1$	
	0 0111 001	0	$9/8 * 1 = 9/8$	closest to 1 above
	0 0111 010	0	$10/8 * 1 = 10/8$	
	...			
	0 1110 110	7	$14/8 * 128 = 224$	
	0 1110 111	7	$15/8 * 128 = 240$	
	0 1111 000	n/a	inf	largest norm

$\hookrightarrow \text{exp} \frac{\text{frac}}{E}$

## Optimization

- Watch innermost loop
- Move code out of the loop
- Multiplication can be replaced with shifts
- Procedure calls are treated as black-boxes

## Architecture

- Superscalar processor: can execute multiple instructions in one cycle
- Cycles per element:  $\cdot \text{Execution time} = \text{CPE} \cdot n + \text{overhead}$

## Caches

- Miss rate: fraction of memory not found in cache
- Hit time: time to deliver a line in the cache to the CPU
- Miss penalty: additional time required because of a miss
  - $\hookrightarrow 99\% \text{ hits is twice as good as } 97\% \text{ hits}$ 
    - Cache hit time of 1 cycle > miss penalty of 100 cycles
    - 97% hits:  $1 \text{ cycle} + 0,03 \cdot 100 \text{ cycles} = 4 \text{ cycles}$
    - 99% hits:  $1 \text{ cycle} + 0,01 \cdot 100 \text{ cycles} = 2 \text{ cycles}$

$\hookrightarrow$  That is, why miss rate is used, instead of hit-rate

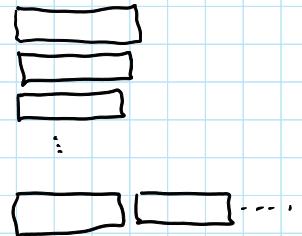
- Cache miss:
  - Cold / Compulsory: first time you access a block
  - Conflict miss: multiple blocks map to same location
  - Capacity miss: Cache too small for all active blocks
  - Coherency miss: multiprocessor misses

- $S$ : # sets
- $E$ : # Blocks/Line per Set
- $B$ : # Bytes per block



1. Set is located with set-index
2. Check if any block/line matches the tag
3. Data starting at the offset

- Direct-mapped-cache:
  - One block per set
  - No associativity
- Fully-associative-cache:
  - Only one set
- Private: only one core sees the cache
- Shared: for multiple cores
- Inclusive: Everything in this cache is also in lower-level-caches
- Exclusive: Everything in this cache isn't in lower-level-caches



### Cache writes

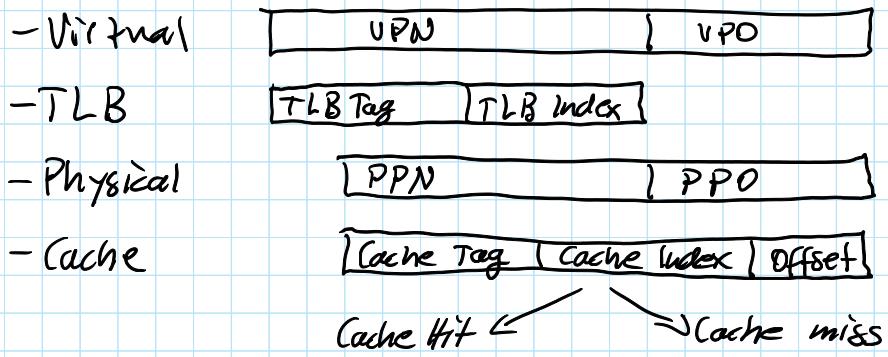
- Write hit
  - Writethrough cache:
    - Write immediately back to memory
  - Writeback cache:
    - Wait to write back, until block gets replaced
    - Needs a dirty bit
- Write miss
  - Write-allocate cache:
    - Load into cache, update block in cache
  - No-write-allocate cache:
    - Writes immediately to memory

## Exceptions

- Each type of exception-event has a unique exception-number =<sup>1</sup> index into exception-table
- Kernel-mode
  - Access to system state
  - Some new instructions / registers
  - For resolving the exception
- Synchronous exceptions
  - As a result of executing an instruction
  - Trap: intentional exception e.g. breakpoint
  - faults: recoverable e.g. page-fault
  - aborts: unrecoverable e.g. errors
- Asynchronous exceptions
  - Events of extern e.g. reset-buttons
  - Also called interrupts

## Virtual memory

- Advantages:
  - Linear address space for each program
  - isolates the programs
  - Efficient use of main-memory (=RAM)
- Page hit:
  1. Process sends virtual address to MMU
  2. MMU fetches physical address from TLB
- Page fault:
  1. Process sends virtual address to MMU
  2. MMU gets page from the disk



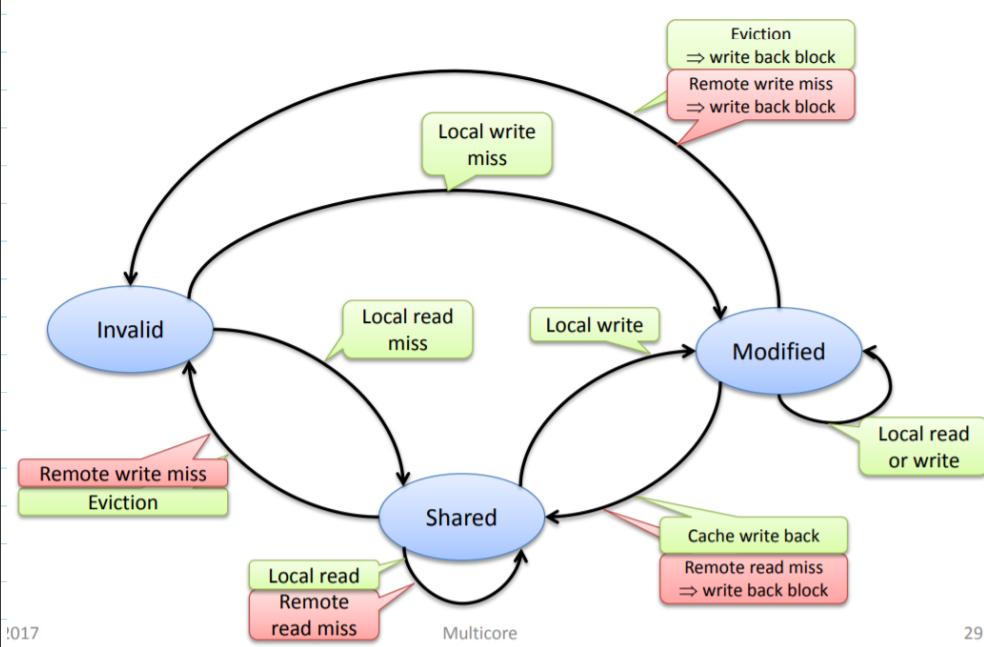
### Cache addressing schemes

- Virtually-indexed, virtually tagged (VV) } Problem, same names for different physical locations
- Virtually-indexed, physically tagged (VP)
- Physically-indexed, physically tagged (PP)
- Physically-indexed, virtually tagged (PV)

### Multiprocessing

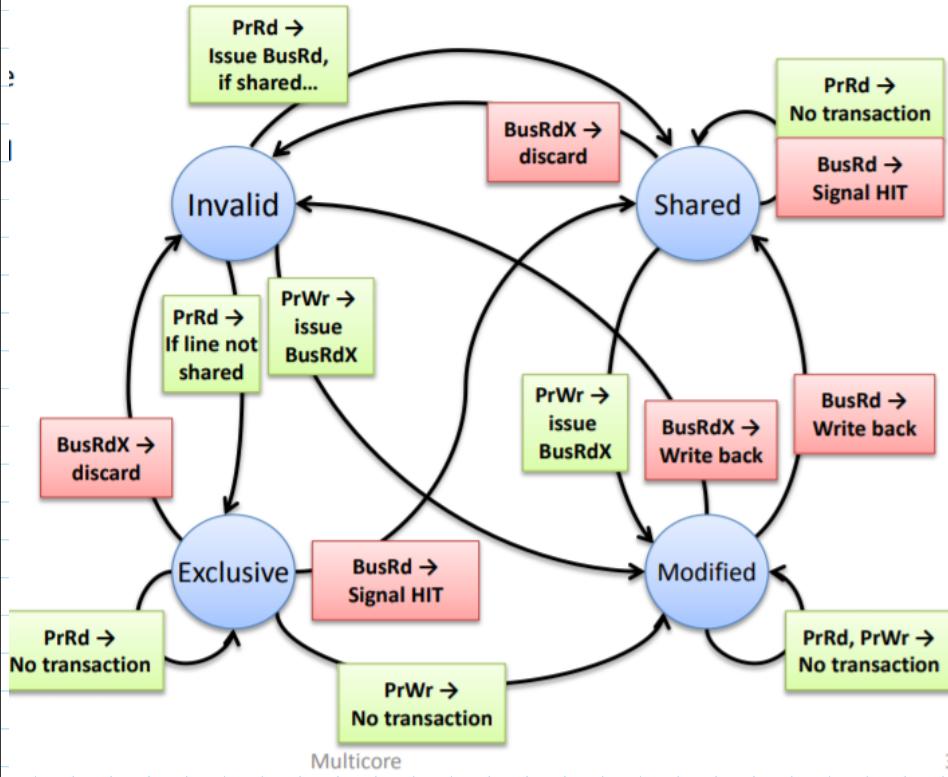
- Coherence: Values in all caches match each other
- Consistency: the order in which changes to the memory are seen by the CPUs

### Cache Coherence MSI



- Line can be dirty in only one cache
- Invalidate: removes element from cache, after writing back
- flush: write back

## Cache Coherence MESI



- Exclusive: Wenn datei niemand anders haben kann

- Advantage: If we are in exclusive and read/write, we must not broadcast that.

34

## Devices

- Registers:
  - Obtain status info
  - Read input data
  - Write output data
- Device registers bypass the cache
- How does the CPU knows, when device is finished => interrupts
- Direct memory access
  - Bypass CPU to transfer data directly between Device > Register
  - Very simple processor => Only copies data
- Too many interrupts
  - Asynchronous buffering - ring
- DEVICE INITIALIZATION
  1. Wait for hardware to initialize
  2. Stop the device from doing anything
  3. Create buffer ring
  4. Write in register to start